



Министерство науки и высшего образования Российской Федерации
Федеральное государственное бюджетное образовательное учреждение высшего
образования
«Московский государственный технический университет имени Н.Э. Баумана
(национальный исследовательский университет)»
(МГТУ им. Н.Э. Баумана)

Факультет «Информатика и системы управления»
Кафедра ИУ5 «Системы обработки информации и управления»

Курс «Методы машинного обучения»

Отчет по лабораторной работе 7

Выполнил:
студент группы ИУ5-24М Поташников
М.Д.

20.05.2023

Лабораторная работа 7. Реализовать алгоритм Actor-Critic для любой среды

```
In [1]: import torch import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import torch.distributions as distributions

import matplotlib.pyplot as plt import
numpy as np
import gym import
tqdm
```

```
In [2]: train_env = gym.make('CartPole-v1') test_env
= gym.make('CartPole-v1')
```

```
In [3]: SEED = 1234

train_env.seed(SEED); test_env.seed(SEED+1);
np.random.seed(SEED);
torch.manual_seed(SEED);
```

```
In [4]: class MLP(nn.Module): def __init__(self, input_dim,
hidden_dim, output_dim): super().__init__()

self.fc_1 = nn.Linear(input_dim, hidden_dim)
self.fc_2 = nn.Linear(hidden_dim, output_dim)
def forward(self,
x): x =
self.fc_1(x) x =
F.relu(x) x =
self.fc_2(x)
return x
```

```
input_dim = train_env.observation_space.shape[0]
hidden_dim = 32
```

In

```
[5]: class ActorCritic(nn.Module):      def
      __init__(self, actor, critic):
      super().__init__()

      self.actor = actor
      self.critic = critic
      def
      forward(self, state):

          action_pred = self.actor(state)
          value_pred = self.critic(state)

          return action_pred, value_pred
```

In

```
[6]:
In
[7]: def init_weights(m):      if type(m) ==
nn.Linear:
torch.nn.init.kaiming_normal_(m.weight)
      m.bias.data.fill_(0)
```

In [8]:

```
device = torch.device('cuda')
```

```
output_dim = train_env.action_space.n
```

```
def calculate_returns(rewards, discount_factor, device, normalize = True):

    returns = []
```

In [9]:

```
def train(env, policy, optimizer, discount_factor, device):
```

```
    policy.train()
```

```
        log_prob_actions = []
```

```
    entropies = []
```

```
    value_preds = []
```

```
    rewards = []    done =
```

```
    False    episode_reward
    = 0
```

```
    state = env.reset()
```

```
    while not
```

```
done:
```

```
        state = torch.FloatTensor(state).unsqueeze(0).to(device)
```

```
        action_pred, value_pred = policy(state)
```

```
        action_prob = F.softmax(action_pred, dim = -1)
```

```
        dist = distributions.Categorical(action_prob)
```

```
        action = dist.sample()
```

```
        log_prob_action = dist.log_prob(action)
```

```
        entropy = dist.entropy()
```

```
        state, reward, done, _ = env.step(action.item())
```

```
        log_prob_actions.append(log_prob_action)
```

```
    entropies.append(entropy)
```

```
        value_preds.append(value_pred.squeeze(0))
```

```
    rewards.append(reward)
```

```
        episode_reward += reward
```

```
    log_prob_actions = torch.cat(log_prob_actions)
```

```
    entropies = torch.cat(entropies)    value_preds =
```

```
    torch.cat(value_preds)
```

```
    returns = calculate_returns(rewards, discount_factor, device)
```

```
    advantages = calculate_advantages(returns, value_preds)
```

```
        loss = update_policy(advantages, log_prob_actions, returns, value_preds, entropies,
```

```
        return loss, episode_reward
```

```
def calculate_advantages(returns, pred_values, normalize = True):
```

```
    advantages = returns - pred_values
```

```
        if
```

```
normalize:
```

```
        advantages = (advantages - advantages.mean()) / advantages.std()
```

```
    return advantages
```

```

In [12]: def evaluate(env, policy, device):

    policy.eval()

    done = False
    episode_reward = 0

    state = env.reset()
    while not
done:

        state = torch.FloatTensor(state).unsqueeze(0).to(device)
        with
torch.no_grad():

            action_pred, _ = policy(state)

            action_prob = F.softmax(action_pred, dim = -1)

In [13]:         action = torch.argmax(action_prob, dim = -1)


        R = 0         for r in
reversed(rewards):         R = r
+ R * discount_factor
returns.insert(0, R)

        returns = torch.tensor(returns).to(device)
        if normalize:         returns = (returns -
returns.mean()) / returns.std()

    return returns

```

```

In
[14]: def update_policy(advantages, log_prob_actions, returns, value_preds, entropies, optimizer

    returns = returns.detach()

    policy_loss = -(advantages * log_prob_actions).mean()
value_loss = F.smooth_l1_loss(returns, value_preds)
        state, reward, done, _ = env.step(action.item())
optimizer.zero_grad()
        episode_reward += reward
loss = policy_loss + value_loss * 0.5 - entropies.mean() * 0.01
        return episode_reward
loss.backward()

optimizer.step()

    return loss.item()

```

de

```
n_runs = 5 max_episodes
= 300 discount_factor =
0.99

train_rewards = torch.zeros(n_runs, max_episodes)
test_rewards = torch.zeros(n_runs, max_episodes) device
= torch.device('cpu')
for run in
u
range(n_runs):
n
:
actor = MLP(input_dim, hidden_dim, output_dim)
critic = MLP(input_dim, hidden_dim, 1) actor_critic
= ActorCritic(actor, critic) actor_critic =
:
actor_critic.to(device)
actor_critic.apply(init_weights)
1
optimizer = optim.Adam(actor_critic.parameters(), lr=1e-2)
0
for episode in tqdm.tqdm(range(max_episodes), desc=f'Run:
0
{run}'):
%
|
| loss, train_reward = train(train_env, actor_critic, optimizer, discount_factor,
|
| test_reward = evaluate(test_env, actor_critic, device)
|
| train_rewards[run][episode] = train_reward
test_rewards[run][episode] = test_reward
|
|
| 300/300 [01:33<00:00, 3.21it/s]
Run: 1: 100%|██████████| 300/300 [01:40<00:00, 2.99it/s]
Run: 2: 100%|██████████| 300/300 [01:29<00:00, 3.35it/s]
Run: 3: 100%|██████████| 300/300 [01:23<00:00, 3.60it/s]
```

```
R
u
n
idxs = range(max_episodes)
fig, ax = plt.subplots(1, figsize=(10,6)) ax.plot(idxs,
test_rewards.mean(0))
ax.fill_between(idxs, test_rewards.min(0).values, test_rewards.max(0).values, alpha=0.1)
ax.set_xlabel('Steps') ax.set_ylabel('Rewards');
1
0
0%|██████████| 300/300 [01:23<00:00, 3.59it/s]
```


In

[15]:

In

[16]:

```
x = torch.randn(2, 10) y =
torch.randn(2, 10)
print(F.smooth_l1_loss(x, y))
print(F.mse_loss(x, y))
```



tensor(0.5806) tensor(1.4047)

This is formatted as code

This is formatted as code